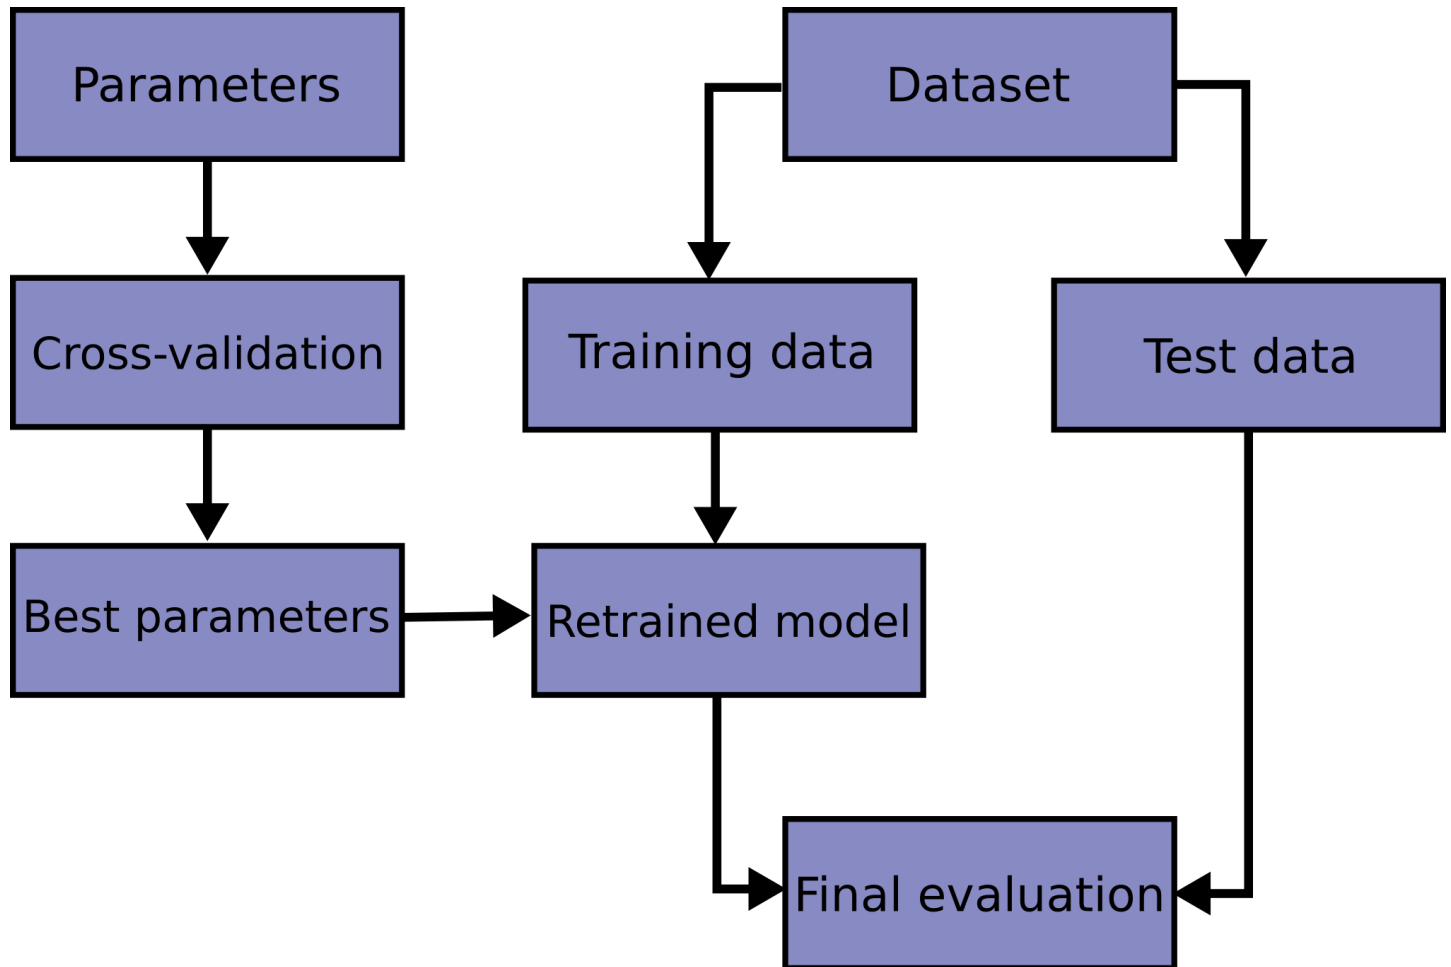


%%[markdown]



```
In [1]: '''
robust versions of logistic regression
support vector machines
random forests
gradient boosted decision trees (XGBoost)
neural networks

Documentacion oficial:
Cross Validation: https://scikit-learn.org/stable/modules/cross_validation.html
GridSearchCV: https://scikit-learn.org/stable/modules/grid_search.html#grid-search
Pipelines: https://scikit-learn.org/stable/modules/compose.html#

https://www.kdnuggets.com/2020/06/simplifying-mixed-feature-type-preprocessing-scikit-learn-pipelines.html
https://www.kaggle.com/kritidoneria/explainable-ai-eli5-lime-and-shap
https://shap.readthedocs.io/en/latest/example_notebooks/tabular_examples/model_agnostic/Diabetes%20regression.html
https://kiwidamien.github.io/introducing-the-column-transformer.html
https://medium.com/analytics-vidhya/shap-part-2-kernel-shap-3c11e7a971b1

*Pendientes*
1)cual tipo de score es el mejor
2)como entrenar con dataset desbalanceado 10:1
3)como evaluar varios modelos adentro de gridsearchcv
'''

import os, sys
import pandas as pd
import numpy as np
import seaborn as sns; sns.set()
from matplotlib import pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from scipy import stats
from sklearn.svm import SVC
from matplotlib.colors import ListedColormap
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import metrics
from sklearn.model_selection import GridSearchCV
import joblib
sns.set(color_codes=True)
from sklearn.pipeline import Pipeline
from sklearn.compose import make_column_selector
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score
```

```
In [2]: path = "/home/nacho/Documents/coronavirus/COVID-19_Paper/"
os.chdir(os.path.join(path))
```

```
In [3]: #%%Valida si existen las carpetas
try:
    os.makedirs("plots")
    os.makedirs("models")
except FileExistsError:
    pass
```

```
In [3]: #%%gridsearchcv
#chechar stratify
#chechar tipo de scoring
#sklearn.metrics.SCORERS.keys()
def gridsearchcv(X_train, X_test, y_train, y_test):
    #####
    # Scale numeric values
    num_transformer = Pipeline(steps=[
        ('scaler', MinMaxScaler())])
    #('scaler', StandardScaler())

    preprocessor = ColumnTransformer(
        remainder='passthrough',
        transformers=[
            ('num', num_transformer, make_column_selector(pattern='EDAD'))
        ])
    #####

    pipe_steps = [
        #('scaler', StandardScaler()),
        ('preprocessor', preprocessor),
        ('SupVM', SVC(kernel='rbf', probability=True)) #agregar kernel linear
    ]

    param_grid= {
        'SupVM__C': [0.1, 0.5, 1, 10, 30, 40, 50, 75, 100, 250, 500, 750, 1000],
        'SupVM__gamma': [0.0001, 0.001, 0.005, 0.01, 0.05, 0.07, 0.1, 0.5, 1, 5, 10, 50, 75, 100]
    }

    pipeline = Pipeline(pipe_steps)
    grid = GridSearchCV(pipeline, param_grid, refit = True, cv = 5, verbose = 3, n_jobs=-1, scoring='balanced_accuracy')# 'f1'
    grid.fit(X_train, y_train)
    print("Best-Fit Parameters From Training Data:\n", grid.best_params_)
    grid_predictions = grid.best_estimator_.predict(X_test)
    report = classification_report(y_test, grid_predictions, output_dict=True)
    report = pd.DataFrame(report).transpose()
    print(report)
    print(confusion_matrix(y_test, grid_predictions))
    return grid, report
```

```
In [3]: ##CASO 1: prediccion de hospitalizacion por covid
hosp_data = pd.read_csv("prediction_data/df_casol.zip")
#Porcentaje de informacion del dataset
hosp_data = hosp_data.sample(frac=0.001)

#separar datos
X = hosp_data.loc[:, hosp_data.columns != 'TIPO_PACIENTE']
y = hosp_data.loc[:, 'TIPO_PACIENTE']
print(y.value_counts())
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42, stratify=y, shuffle=True)

0    1697
1     400
Name: TIPO_PACIENTE, dtype: int64
```

```
In [10]: ##Entrenamiento (opcional)
#--->train
grid, grid_report= gridsearchcv(X_train, X_test, y_train, y_test)
#guarda el modelo y su reporte
joblib.dump(grid, 'models/hosp_data_grid.pkl', compress = 1)
grid_report.to_csv("models/hosp_data_grid_report.csv", index=True)
'''

Best-Fit Parameters From Training Data:
{'SupVM_C': 500, 'SupVM_gamma': 1}
precision    recall    f1-score    support
0      0.845872   0.964435   0.901271   478.00000
1      0.645833   0.269565   0.380368   115.00000
accuracy      0.829680   0.829680   0.829680   0.82968
macro avg     0.745852   0.617000   0.640819   593.00000
weighted avg   0.807078   0.829680   0.800253   593.00000
[[461  17]
 [ 84  31]]
'''

Fitting 5 folds for each of 154 candidates, totalling 770 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 16 tasks | elapsed: 2.6s
[Parallel(n_jobs=-1)]: Done 112 tasks | elapsed: 6.6s
[Parallel(n_jobs=-1)]: Done 272 tasks | elapsed: 15.4s
[Parallel(n_jobs=-1)]: Done 496 tasks | elapsed: 43.4s
[Parallel(n_jobs=-1)]: Done 770 out of 770 | elapsed: 3.2min finished
Best-Fit Parameters From Training Data:
{'SupVM_C': 75, 'SupVM_gamma': 5}
precision    recall    f1-score    support
0      0.828897   0.923729   0.873747   472.000000
1      0.462687   0.256198   0.329787   121.000000
accuracy      0.787521   0.787521   0.787521   0.787521
macro avg     0.645792   0.589964   0.601767   593.000000
weighted avg   0.754173   0.787521   0.762754   593.000000
[[436  36]
 [ 90  31]]
```

```
Out[10]: "\nBest-Fit Parameters From Training Data:\n {'SupVM_C': 500, 'SupVM_gamma': 1}\n
precision    recall    f1-score    support
\n0      0.845872   0.964435   0.901271   478.00000\n1      0.645833   0.269565   0.380368   115.00000\naccuracy      0.829680   0.829680   0.829680   0.82968\nmacro avg     0.745852   0.617000   0.640819   593.00000\nweighted avg   0.807078   0.829680   0.800253   593.00000\n[[461  17]\n [ 84  31]]\n"
```

```
In [4]: ##importa el modelo y su rendimiento
grid = joblib.load('models/hosp_data_grid.pkl')
grid_report = pd.read_csv("models/hosp_data_grid_report.csv", index_col=0)
print("best score from grid search: %f" % grid.best_estimator_.score(X_test, y_test))

best score from grid search: 0.809524
```

```
In [11]: #prediccion
grid.predict(pd.DataFrame(X_train.iloc[20, :].values.reshape(1, -1), columns = X_train.columns))
```

```
Out[11]: array([0])
```

```
In [12]: #prediccion con probabilidad
grid.predict_proba(pd.DataFrame(X_train.iloc[20, :].values.reshape(1, -1), columns = X_train.columns))
```

```
Out[12]: array([[0.71358794, 0.28641206]])
```

```
In [6]: #metodos y atributos del objeto grid
dir(grid)
```

```
Out[6]: ['_abstractmethods_',
'_class_',
'_delattr_',
'_dict_',
'_dir_',
'_doc_',
'_eq_',
'_format_',
'_ge_',
'_getattr_',
'_getstate_',
'_gt_',
'_hash_',
'_init_',
'_init_subclass_',
'_le_',
'_lt_',
'_module_',
'_ne_',
'_new_',
'_reduce_',
'_reduce_ex_',
'_repr_',
'_setattr_',
'_setstate_',
'_sizeof_',
'_str_',
'_subclasshook_',
'_weakref_',
'_abc_cache',
'_abc_negative_cache',
'_abc_negative_cache_version',
'_abc_registry',
'_check_is_fitted',
'_estimator_type',
'_format_results',
'_get_param_names',
'_get_tags',
'_more_tags',
'_pairwise',
'_required_parameters',
'_run_search',
'_best_estimator_',
'_best_index_',
'_best_params_',
'_best_score_',
'_classes_',
'_cv',
'_cv_results_',
'_decision_function',
'_error_score',
'_estimator',
'_fit',
'_get_params',
'_iid',
'_inverse_transform',
'_multimetric_',
'_n_iter',
'_n_jobs',
'_n_splits',
'_param_distributions',
'_pre_dispatch',
'_predict',
'_predict_log_proba',
'_predict_proba',
'_random_state',
'_refit',
'_refit_time_',
'_return_train_score',
'_score',
'_scorer_',
'_scoring',
'_set_params',
'_transform',
'_verbose']
```

```
In [13]: grid.best_score_
```

```
Out[13]: 0.3233383946258245
```

```
In [15]: grid.best_params_
```

```
Out[15]: {'clf__selected_model': ('gb',
{'learning_rate': 0.15,
 'max_depth': 9,
 'max_features': 'sqrt',
 'n_estimators': 50,
 'subsample': 0.618})}
```

```
In [14]: grid.best_estimator_
```

```
Out[14]: Pipeline(memory=None,
  steps=[('preprocessor',
    ColumnTransformer(n_jobs=None, remainder='passthrough',
      sparse_threshold=0.3,
      transformer_weights=None,
      transformers=[('num',
        Pipeline(memory=None,
          steps=[('scaler',
            MinMaxScaler(copy=True,
              feature_range=(0,
                1))]),
          verbose=False),
        <sklearn.compose._column_transformer.make_column_selector object at 0x7fe629f43be0>)],
        ver...
          loss='deviance',
          max_depth=9,
          max_features='sqrt',
          max_leaf_nodes=None,
          min_impurity_decrease=0.0,
          min_impurity_split=None,
          min_samples_leaf=1,
          min_samples_split=2,
          min_weight_fraction_leaf=0.0,
          n_estimators=50,
          n_iter_no_change=None,
          presort='deprecated',
          random_state=None,
          subsample=0.618,
          tol=0.0001,
          validation_fraction=0.1,
          verbose=0,
          warm_start=False))),
    verbose=False)
```

```
In [7]: df_grid.columns.tolist()
```

```
Out[7]: ['mean_fit_time',
        'std_fit_time',
        'mean_score_time',
        'std_score_time',
        'param_clf_selected_model',
        'params',
        'split0_test_f1',
        'split1_test_f1',
        'split2_test_f1',
        'split3_test_f1',
        'split4_test_f1',
        'split5_test_f1',
        'split6_test_f1',
        'split7_test_f1',
        'split8_test_f1',
        'split9_test_f1',
        'split10_test_f1',
        'split11_test_f1',
        'split12_test_f1',
        'split13_test_f1',
        'split14_test_f1',
        'split15_test_f1',
        'split16_test_f1',
        'split17_test_f1',
        'split18_test_f1',
        'split19_test_f1',
        'split20_test_f1',
        'split21_test_f1',
        'split22_test_f1',
        'split23_test_f1',
        'split24_test_f1',
        'split25_test_f1',
        'split26_test_f1',
        'split27_test_f1',
        'split28_test_f1',
        'split29_test_f1',
        'split30_test_f1',
        'split31_test_f1',
        'split32_test_f1',
        'split33_test_f1',
        'split34_test_f1',
        'split35_test_f1',
        'split36_test_f1',
        'split37_test_f1',
        'split38_test_f1',
        'split39_test_f1',
        'split40_test_f1',
        'split41_test_f1',
        'split42_test_f1',
        'split43_test_f1',
        'split44_test_f1',
        'split45_test_f1',
        'split46_test_f1',
        'split47_test_f1',
        'split48_test_f1',
        'split49_test_f1',
        'mean_test_f1',
        'std_test_f1',
        'rank_test_f1',
        'split0_test_balanced_accuracy',
        'split1_test_balanced_accuracy',
        'split2_test_balanced_accuracy',
        'split3_test_balanced_accuracy',
        'split4_test_balanced_accuracy',
        'split5_test_balanced_accuracy',
        'split6_test_balanced_accuracy',
        'split7_test_balanced_accuracy',
        'split8_test_balanced_accuracy',
        'split9_test_balanced_accuracy',
        'split10_test_balanced_accuracy',
        'split11_test_balanced_accuracy',
        'split12_test_balanced_accuracy',
        'split13_test_balanced_accuracy',
        'split14_test_balanced_accuracy',
        'split15_test_balanced_accuracy',
        'split16_test_balanced_accuracy',
        'split17_test_balanced_accuracy',
        'split18_test_balanced_accuracy',
        'split19_test_balanced_accuracy',
        'split20_test_balanced_accuracy',
        'split21_test_balanced_accuracy',
        'split22_test_balanced_accuracy',
        'split23_test_balanced_accuracy',
        'split24_test_balanced_accuracy',
        'split25_test_balanced_accuracy',
        'split26_test_balanced_accuracy',
        'split27_test_balanced_accuracy',
        'split28_test_balanced_accuracy',
        'split29_test_balanced_accuracy',
        'split30_test_balanced_accuracy',
        'split31_test_balanced_accuracy',
        'split32_test_balanced_accuracy',
        'split33_test_balanced_accuracy',
        'split34_test_balanced_accuracy',
        'split35_test_balanced_accuracy',
        'split36_test_balanced_accuracy',
        'split37_test_balanced_accuracy',
        'split38_test_balanced_accuracy',
        'split39_test_balanced_accuracy',
        'split40_test_balanced_accuracy',
        'split41_test_balanced_accuracy',
        'split42_test_balanced_accuracy',
```

```
'split43_test_balanced_accuracy',
'split44_test_balanced_accuracy',
'split45_test_balanced_accuracy',
'split46_test_balanced_accuracy',
'split47_test_balanced_accuracy',
'split48_test_balanced_accuracy',
'split49_test_balanced_accuracy',
'mean_test_balanced_accuracy',
'std_test_balanced_accuracy',
'rank_test_balanced_accuracy']
```

```
In [9]: df_grid = pd.DataFrame(grid.cv_results_)
df_grid.sort_values(by = ['mean_test_balanced_accuracy'], ascending=False)
```

```
Out[9]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_clf_selected_model	params	split0_test_f1	split1_test_f1	split2_test_f1	split3_test_f1	...	split43_test_f1
0	0.643435	0.056959	0.013391	0.003793	(gb, {'learning_rate': 0.15, 'max_depth': 9, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.318182	0.384615	0.333333	0.186047	...	
9	1.728619	0.153196	0.013713	0.002805	(gb, {'learning_rate': 0.025, 'max_depth': 8, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.315789	0.238095	0.390244	0.108108	...	
1	16.933959	1.364132	0.034127	0.006654	(gb, {'learning_rate': 0.075, 'max_depth': 50, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.297872	0.313725	0.425532	0.186047	...	
4	6.832211	0.269802	0.021366	0.003339	(gb, {'learning_rate': 0.05, 'max_depth': 8, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.285714	0.280000	0.434783	0.232558	...	
2	1.817193	0.063662	0.014718	0.001571	(gb, {'learning_rate': 0.2, 'max_depth': 3, 'm...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.292683	0.243902	0.292683	0.190476	...	
8	4.623009	0.257587	0.018732	0.003551	(gb, {'learning_rate': 0.075, 'max_depth': 8, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.243902	0.285714	0.375000	0.222222	...	
7	18.606113	0.633623	0.029817	0.006988	(gb, {'learning_rate': 0.01, 'max_depth': 50, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.272727	0.346154	0.434783	0.227273	...	
6	0.289662	0.025178	0.012307	0.001844	(gb, {'learning_rate': 0.025, 'max_depth': 3, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.000000	0.303030	0.285714	0.277778	...	
5	12.370118	0.425477	0.025213	0.002602	(gb, {'learning_rate': 0.001, 'max_depth': 10, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.000000	0.071429	0.129032	0.000000	...	
3	0.534879	0.040907	0.012240	0.001531	(gb, {'learning_rate': 0.001, 'max_depth': 8, '...	{'clf__selected_model': 'gb', 'learning_rate': ...	0.000000	0.000000	0.000000	0.000000	...	

10 rows × 112 columns

```
In [37]: #%%Pipeline without gridsearchcv

num_transformer = Pipeline(steps=[
    ('scaler', MinMaxScaler())])

preprocessor = ColumnTransformer(
    remainder='passthrough',
    transformers=[
        ('num', num_transformer, ['EDAD'])])

svc_model = Pipeline([
    #('scaler', StandardScaler()),
    ('preprocessor', preprocessor),
    ('SupVM', SVC(kernel='rbf', probability=True))])

#svc_model.set_params(**grid.best_params_)
svc_model.set_params(**{'SupVM_C': 500, 'SupVM_gamma': 1})
#svc_model.set_params(**{'SupVM_C': 75, 'SupVM_gamma': 5})
#print('Model params: ', svc_model.get_params("model"))

svc_model.fit(X_train, y_train)
y_pred = svc_model.predict(X_test)
metrics.accuracy_score(y_test, y_pred)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.93	0.88	478
1	0.48	0.28	0.35	115
accuracy			0.80	593
macro avg	0.66	0.60	0.62	593
weighted avg	0.77	0.80	0.78	593

```
In [39]: #train evaluation
model_score = cross_val_score(svc_model, X_train, y_train, cv=StratifiedKFold(n_splits=5))
print("5 split scores:" ,model_score)
print("Mean of the 5 split scores:" ,model_score.mean())

5 split scores: [0.8340249 0.81327801 0.80912863 0.825 0.82916667]
Mean of the 5 split scores: 0.8221196403872753
```

```
In [36]: #final evaluation
final_score = cross_val_score(svc_model, X_test, y_test, cv=StratifiedKFold(n_splits=5))
print("5 split scores:" ,final_score)
print("Mean of the 5 split scores:" ,final_score.mean())

5 split scores: [0.79831933 0.79831933 0.8487395 0.78813559 0.80508475]
Mean of the 5 split scores: 0.8077196980487111
```

```
In [30]: #%%se preprocesan los datos
X_train_scaled = pd.DataFrame(svc_model.named_steps['preprocessor'].fit_transform(X_train), columns = X_train.columns)
X_test_scaled = pd.DataFrame(svc_model.named_steps['preprocessor'].fit_transform(X_test), columns = X_test.columns)
#cambia el orden de edad y sexo
X_train_scaled[['EDAD', 'SEXO']] = X_train_scaled[['SEXO', 'EDAD']]
X_test_scaled[['EDAD', 'SEXO']] = X_test_scaled[['SEXO', 'EDAD']]
```



```
In [31]: X_train_scaled.head()
```

Out[31]:

	SEXO	EDAD	EMBARAZO	DIABETES	EPOC	ASMA	INMUSUPR	HIPERTENSION	CARDIOVASCULAR	OBEIDAD	RENAL_CRONICA	TABAQUISMO
0	1.0	0.271739	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.521739	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
2	1.0	0.500000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	1.0	0.543478	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	1.0	0.260870	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

```
In [33]: X_test_scaled.head()
```

Out[33]:

	SEXO	EDAD	EMBARAZO	DIABETES	EPOC	ASMA	INMUSUPR	HIPERTENSION	CARDIOVASCULAR	OBEIDAD	RENAL_CRONICA	TABAQUISMO
0	0.0	0.623656	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
1	1.0	0.301075	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.505376	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.462366	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	1.0	0.236559	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0

```
In [18]: y_test.iloc[20]
```

Out[18]: 1

```
In [17]: X_test.iloc[20]
```

Out[17]: SEXO 0
EDAD 26
EMBARAZO 0
DIABETES 0
EPOC 0
ASMA 0
INMUSUPR 0
HIPERTENSION 0
CARDIOVASCULAR 0
OBEIDAD 0
RENAL_CRONICA 0
TABAQUISMO 0
Name: 932525, dtype: int64

```
In [19]: ##prueba el modelo con X_test  
svc_model.named_steps['SupVM'].predict(X_test_scaled.iloc[20, :].values.reshape(1, -1))
```

Out[19]: array([0])

```
In [20]: svc_model.named_steps['SupVM'].predict_proba(X_test_scaled.iloc[20, :].values.reshape(1, -1))
```

Out[20]: array([[0.91288434, 0.08711566]])

```
In [21]: import shap  
# use Kernel SHAP to explain test set predictions  
explainer = shap.KernelExplainer(model = svc_model.named_steps['SupVM'].predict_proba, data = X_train_scaled, link = 'logit')  
shap_values = explainer.shap_values(X = X_test_scaled, nsamples = 30, ll_reg="num_features(12)")
```

Using 1203 background data samples could cause slower run times. Consider using shap.sample(data, K) or shap.kmeans(data, K) to summarize the background as K samples.
100%|██████████| 593/593 [20:32<00:00, 2.08s/it]


```
In [22]: print(f'length of SHAP values: {len(shap_values)}')  
print(f'Shape of each element: {shap_values[0].shape}')
```

length of SHAP values: 2
Shape of each element: (593, 12)


```
In [23]: #prediction and probability of model  
print(f'Prediction for 1st sample in X_test: ', svc_model.named_steps['SupVM'].predict(X_test_scaled.iloc[[0], :])[0])  
print(f'Prediction probability for 1st sample in X_test: ', svc_model.named_steps['SupVM'].predict_proba(X_test_scaled.iloc[[0], :])[0])
```

Prediction for 1st sample in X_test: 0
Prediction probability for 1st sample in X_test: [0.91075977 0.08924023]

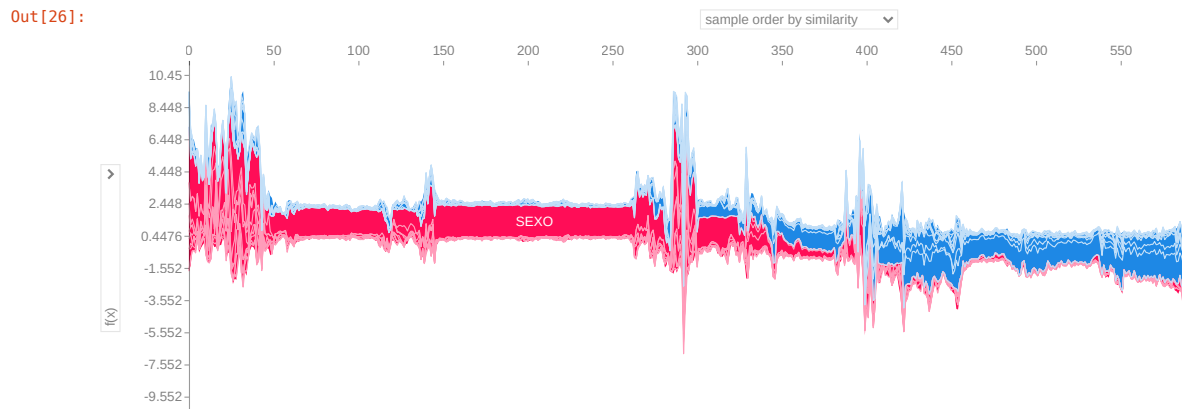
```
In [24]: # plot the SHAP values for the false (0) output of the first instance  
shap.initjs()  
shap.force_plot(explainer.expected_value[0], shap_values[0][0, :], X_test_scaled.iloc[0, :], link="logit")
```



```
In [25]: shap.initjs()
shap.force_plot(explainer.expected_value[1], shap_values[1][0,:], X_test_scaled.iloc[0,:], link="logit")
```






```
In [26]: #Explaining the Prediction for all samples in Test set
#no hospitalizado
shap.initjs()
shap.force_plot(explainer.expected_value[0], shap_values[0], X_test_scaled)
```



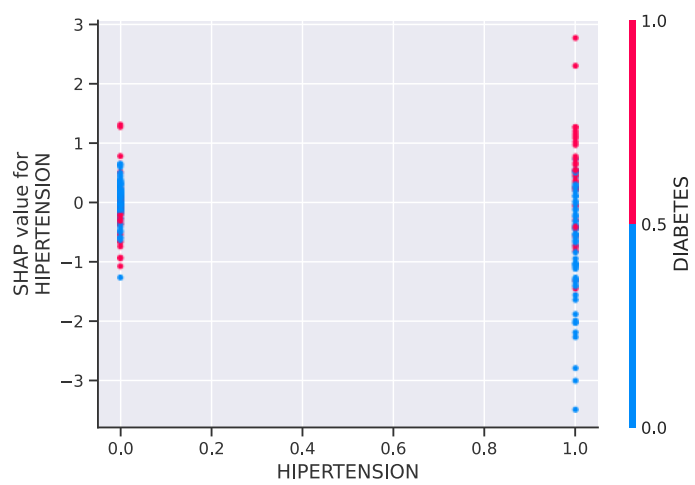
```
In [27]: #si hospitalizado
shap.initjs()
shap.force_plot(explainer.expected_value[1], shap_values[1], X_test_scaled)
```



```
In [28]: #SHAP Summary Plots
shap.summary_plot(shap_values, X_test)
shap.summary_plot(shap_values[1], X_test_scaled)
shap.summary_plot(shap_values[0], X_test_scaled)
```



```
In [29]: #SHAP Dependence Plots
shap.dependence_plot("HIPERTENSION", shap_values[1], X_test_scaled)
```



```
In [ ]:
```